

20bsit154-ass-2

March 30, 2023

```
[1]: #1.      Binning using Python
      #1.      Implement Binning using cut and qcut methods.
      #2.      Also, transform the bins values.

import pandas as pd
na1 = ["n.a.", "not available"]

df = pd.read_csv("D:/sem6/dma/practical/Automobile_data.csv", na_values = na1)
```

```
[2]: print(df)
```

	index	company	body-style	wheel-base	length	engine-type \
0	0	alfa-romero	convertible	88.6	168.8	dohc
1	1	alfa-romero	convertible	88.6	168.8	dohc
2	2	alfa-romero	hatchback	94.5	171.2	ohcv
3	3	audi	sedan	99.8	176.6	ohc
4	4	audi	sedan	99.4	176.6	ohc
..
56	81	volkswagen	sedan	97.3	171.7	ohc
57	82	volkswagen	sedan	97.3	171.7	ohc
58	86	volkswagen	sedan	97.3	171.7	ohc
59	87	volvo	sedan	104.3	188.8	ohc
60	88	volvo	wagon	104.3	188.8	ohc

	num-of-cylinders	horsepower	average-mileage	price
0	four	111	21	13495.0
1	four	111	21	16500.0
2	six	154	19	16500.0
3	four	102	24	13950.0
4	five	115	18	17450.0
..
56	four	85	27	7975.0
57	four	52	37	7995.0
58	four	100	26	9995.0
59	four	114	23	12940.0
60	four	114	23	13415.0

[61 rows x 10 columns]

```
[43]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 9 entries, 0 to 8
Data columns (total 10 columns):
 #   Column          Non-Null Count  Dtype
---  -
 0   index           9 non-null      int64
 1   company         9 non-null      object
 2   body-style      9 non-null      object
 3   wheel-base      9 non-null      float64
 4   length          9 non-null      float64
 5   engine-type     9 non-null      object
 6   num-of-cylinders 9 non-null      object
 7   horsepower      8 non-null      float64
 8   average-mileage 9 non-null      int64
 9   price           7 non-null      float64
dtypes: float64(4), int64(2), object(4)
memory usage: 848.0+ bytes
```

```
[4]: df['wheel-base'].min()
```

```
[4]: 88.4
```

```
[6]: df['wheel-base'].max()
```

```
[6]: 120.9
```

```
[11]: #cut
bins = [88,100,130]
labels = ['small','medium']
```

```
[12]: df['grade'] = pd.cut(df['wheel-base'], bins =bins, labels = labels,
↪include_lowest=True)
```

```
[13]: print(df['grade'], df['wheel-base'])
```

```
0    small
1    small
2    small
3    small
4    small
...
56   small
57   small
58   small
59  medium
60  medium
```

```

Name: grade, Length: 61, dtype: category
Categories (2, object): ['small' < 'medium'] 0      88.6
1      88.6
2      94.5
3      99.8
4      99.4
...
56     97.3
57     97.3
58     97.3
59    104.3
60    104.3
Name: wheel-base, Length: 61, dtype: float64

```

```
[14]: df['grade'].value_counts()
```

```

[14]: small      42
      medium     19
      Name: grade, dtype: int64

```

```

[52]: #qcut

df.describe()

```

```

[52]:
      count      index  wheel-base      length  horsepower  average-mileage  \
count    9.000000     9.000000     9.000000     8.000000         9.000000
mean     4.444444    97.655556   176.177778   113.125000        20.777778
std       3.431877     5.891331     7.150311    17.365298         2.166667
min       0.000000    88.600000   168.800000   101.000000        18.000000
25%       2.000000    94.500000   171.200000   101.750000        19.000000
50%       4.000000    99.800000   176.600000   110.500000        21.000000
75%       6.000000   101.200000   176.800000   112.000000        23.000000
max      10.000000   105.800000   192.700000   154.000000        24.000000

      price
count     7.000000
mean    16667.857143
std     1486.028985
min     13950.000000
25%     16465.000000
50%     16500.000000
75%     17187.500000
max     18920.000000

```

```
[15]: df['grades'] = pd.qcut(df['wheel-base'], q=2)
```

```
[16]: print(df['grades'])
```

```

0      (88.399, 96.3]
1      (88.399, 96.3]
2      (88.399, 96.3]
3      (96.3, 120.9]
4      (96.3, 120.9]
...
56     (96.3, 120.9]
57     (96.3, 120.9]
58     (96.3, 120.9]
59     (96.3, 120.9]
60     (96.3, 120.9]
Name: grades, Length: 61, dtype: category
Categories (2, interval[float64, right]): [(88.399, 96.3] < (96.3, 120.9]]

```

```
[17]: df['grades'].value_counts()
```

```

[17]: (88.399, 96.3]      31
      (96.3, 120.9]      30
      Name: grades, dtype: int64

```

```

[18]: labels = ['small', 'medium']
      df['grades'] = pd.qcut(df['wheel-base'], q=2, labels = labels)
      df['grades'].value_counts()
      print(df['grades'], df['wheel-base'])

```

```

0      small
1      small
2      small
3      medium
4      medium
...
56     medium
57     medium
58     medium
59     medium
60     medium
Name: grades, Length: 61, dtype: category
Categories (2, object): ['small' < 'medium'] 0      88.6
1      88.6
2      94.5
3      99.8
4      99.4
...
56     97.3
57     97.3
58     97.3
59     104.3
60     104.3

```

Name: wheel-base, Length: 61, dtype: float64

```
[19]: df.groupby('grades')['wheel-base'].transform('mean')
```

```
[19]: 0      93.729032
      1      93.729032
      2      93.729032
      3     103.393333
      4     103.393333
```

...

```
56     103.393333
57     103.393333
58     103.393333
59     103.393333
60     103.393333
```

Name: wheel-base, Length: 61, dtype: float64

```
[20]: df.groupby('grades')['wheel-base'].transform('median')
```

```
[20]: 0      94.5
      1      94.5
      2      94.5
      3     101.6
      4     101.6
```

...

```
56     101.6
57     101.6
58     101.6
59     101.6
60     101.6
```

Name: wheel-base, Length: 61, dtype: float64

```
[21]: #2.      Outlier detection and removal
      #1.      Detect the outlier using visualization method
      #2.      Detect the outlier using statistical method
      #3.      Treat the outliers

na1 = ["n.a.", "not available"]

df = pd.read_csv("D:/sem6/dma/practical/Automobile_data.csv", na_values = na1)
```

```
[22]: print(df)
```

	index	company	body-style	wheel-base	length	engine-type	\
0	0	alfa-romero	convertible	88.6	168.8	dohc	
1	1	alfa-romero	convertible	88.6	168.8	dohc	
2	2	alfa-romero	hatchback	94.5	171.2	ohcv	
3	3	audi	sedan	99.8	176.6	ohc	


4	4	audi	sedan	99.4	176.6	ohc
..	
56	81	volkswagen	sedan	97.3	171.7	ohc
57	82	volkswagen	sedan	97.3	171.7	ohc
58	86	volkswagen	sedan	97.3	171.7	ohc
59	87	volvo	sedan	104.3	188.8	ohc
60	88	volvo	wagon	104.3	188.8	ohc

	num-of-cylinders	horsepower	average-mileage	price
0	four	111	21	13495.0
1	four	111	21	16500.0
2	six	154	19	16500.0
3	four	102	24	13950.0
4	five	115	18	17450.0
..
56	four	85	27	7975.0
57	four	52	37	7995.0
58	four	100	26	9995.0
59	four	114	23	12940.0
60	four	114	23	13415.0

[61 rows x 10 columns]

```
[23]: import seaborn as sns
sns.boxplot(x=df['price'])
```


```
[23]: <AxesSubplot:xlabel='price'>
```



```
[24]: # Scatter plot
import matplotlib.pyplot as plt
fig, ax = plt.subplots(figsize = (18,10))
ax.scatter(df['price'], df['num-of-cylinders'])

# x-axis label
ax.set_xlabel('X')


# y-axis label
ax.set_ylabel('Y')
plt.show()
```



```
[25]: # Scatter plot
import matplotlib.pyplot as plt
fig, ax = plt.subplots(figsize = (18,10))
ax.scatter(df['price'], df['company'])

# x-axis label
ax.set_xlabel('X')


# y-axis label
ax.set_ylabel('Y')
plt.show()
```

```
[26]: # Scatter plot
import matplotlib.pyplot as plt
fig, ax = plt.subplots(figsize = (18,10))
ax.scatter(df['price'], df['body-style'])

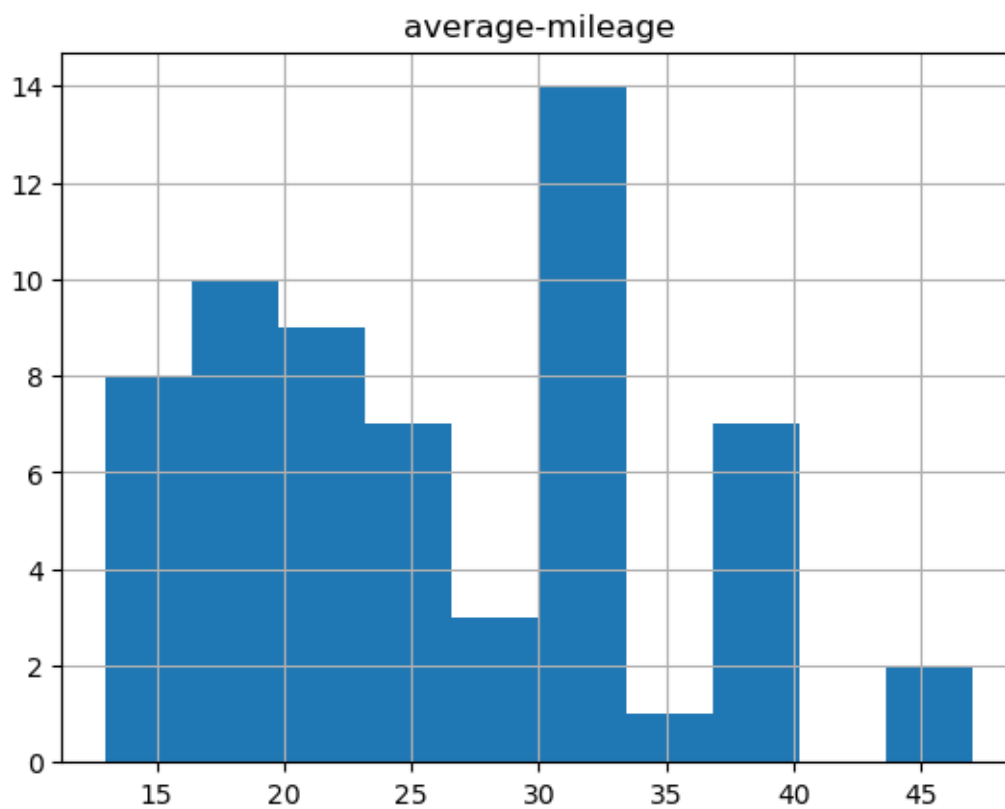
# x-axis label
ax.set_xlabel('X')

# y-axis label
ax.set_ylabel('Y')
plt.show()
```



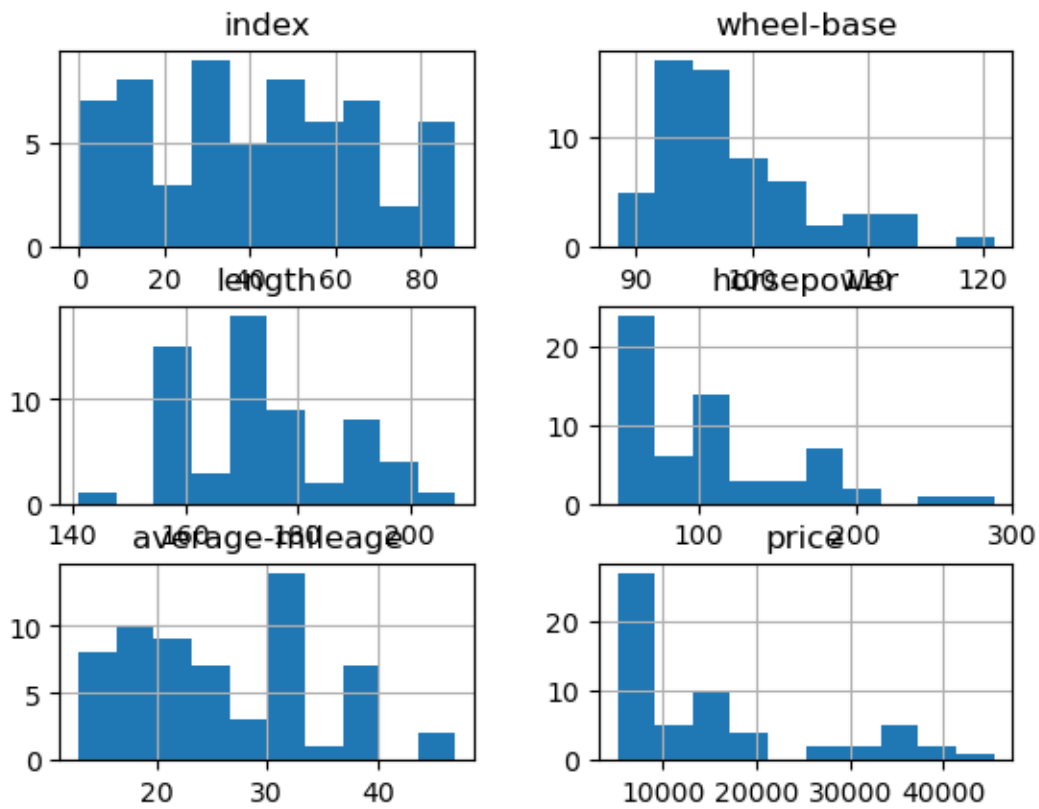
```
[27]: df.hist(column='average-mileage', bins=10)
```

```
[27]: array([[<AxesSubplot:title={'center':'average-mileage'}>]], dtype=object)
```



```
[28]: df.hist()
```

```
[28]: array([[<AxesSubplot:title={'center':'index'}>,
<AxesSubplot:title={'center':'wheel-base'}>],
[<AxesSubplot:title={'center':'length'}>,
<AxesSubplot:title={'center':'horsepower'}>],
[<AxesSubplot:title={'center':'average-mileage'}>,
<AxesSubplot:title={'center':'price'}>]], dtype=object)
```



```
[29]: # Z score
from scipy import stats
import numpy as np

df['price z scores'] = stats.zscore(df['price'], nan_policy='omit')
print(df.head(40))
```

	index	company	body-style	wheel-base	length	engine-type	\
0	0	alfa-romero	convertible	88.6	168.8	dohc	
1	1	alfa-romero	convertible	88.6	168.8	dohc	

2	2	alfa-romero	hatchback	94.5	171.2	ohcv
3	3	audi	sedan	99.8	176.6	ohc
4	4	audi	sedan	99.4	176.6	ohc
5	5	audi	sedan	99.8	177.3	ohc
6	6	audi	wagon	105.8	192.7	ohc
7	9	bmw	sedan	101.2	176.8	ohc
8	10	bmw	sedan	101.2	176.8	ohc
9	11	bmw	sedan	101.2	176.8	ohc
10	13	bmw	sedan	103.5	189.0	ohc
11	14	bmw	sedan	103.5	193.8	ohc
12	15	bmw	sedan	110.0	197.0	ohc
13	16	chevrolet	hatchback	88.4	141.1	l
14	17	chevrolet	hatchback	94.5	155.9	ohc
15	18	chevrolet	sedan	94.5	158.8	ohc
16	19	dodge	hatchback	93.7	157.3	ohc
17	20	dodge	hatchback	93.7	157.3	ohc
18	27	honda	wagon	96.5	157.1	ohc
19	28	honda	sedan	96.5	175.4	ohc
20	29	honda	sedan	96.5	169.1	ohc
21	30	isuzu	sedan	94.3	170.7	ohc
22	31	isuzu	sedan	94.5	155.9	ohc
23	32	isuzu	sedan	94.5	155.9	ohc
24	33	jaguar	sedan	113.0	199.6	dohc
25	34	jaguar	sedan	113.0	199.6	dohc
26	35	jaguar	sedan	102.0	191.7	ohcv
27	36	mazda	hatchback	93.1	159.1	ohc
28	37	mazda	hatchback	93.1	159.1	ohc
29	38	mazda	hatchback	93.1	159.1	ohc
30	39	mazda	hatchback	95.3	169.0	rotor
31	43	mazda	sedan	104.9	175.0	ohc
32	44	mercedes-benz	sedan	110.0	190.9	ohc
33	45	mercedes-benz	wagon	110.0	190.9	ohc
34	46	mercedes-benz	sedan	120.9	208.1	ohcv
35	47	mercedes-benz	hardtop	112.0	199.2	ohcv
36	49	mitsubishi	hatchback	93.7	157.3	ohc
37	50	mitsubishi	hatchback	93.7	157.3	ohc
38	51	mitsubishi	sedan	96.3	172.4	ohc
39	52	mitsubishi	sedan	96.3	172.4	ohc

	num-of-cylinders	horsepower	average-mileage	price	price z scores
0	four	111	21	13495.0	-0.168594
1	four	111	21	16500.0	0.099178
2	six	154	19	16500.0	0.099178
3	four	102	24	13950.0	-0.128049
4	five	115	18	17450.0	0.183831
5	five	110	19	15250.0	-0.012208
6	five	110	19	18920.0	0.314821
7	four	101	23	16430.0	0.092940

8	four	101	23	16925.0	0.137049
9	six	121	21	20970.0	0.497494
10	six	182	16	30760.0	1.369868
11	six	182	16	41315.0	2.310411
12	six	182	15	36880.0	1.915214
13	three	48	47	5151.0	-0.912117
14	four	70	38	6295.0	-0.810176
15	four	70	38	6575.0	-0.785226
16	four	68	31	6377.0	-0.802870
17	four	68	31	6229.0	-0.816058
18	four	76	30	7295.0	-0.721068
19	four	101	24	12945.0	-0.217603
20	four	100	25	10345.0	-0.449286
21	four	78	24	6785.0	-0.766513
22	four	70	38	NaN	NaN
23	four	70	38	NaN	NaN
24	six	176	15	32250.0	1.502640
25	six	176	15	35550.0	1.796699
26	twelve	262	13	36000.0	1.836798
27	four	68	30	5195.0	-0.908196
28	four	68	31	6095.0	-0.827998
29	four	68	31	6795.0	-0.765622
30	two	101	17	11845.0	-0.315623
31	four	72	31	18344.0	0.263494
32	five	123	22	25552.0	0.905790
33	five	123	22	28248.0	1.146027
34	eight	184	14	40960.0	2.278777
35	eight	184	14	45400.0	2.674420
36	four	68	37	5389.0	-0.890909
37	four	68	31	6189.0	-0.819622
38	four	88	25	6989.0	-0.748335
39	four	88	25	8189.0	-0.641405

```
[30]: threshold = 3
      # Position of the outlier
      print(np.where(df['price z scores'] > 2))
```

```
(array([11, 34, 35], dtype=int64),)
```

```
[32]: # IQR
      Q1 = df.quantile(0.25)
      Q3 = df.quantile(0.75)
      IQR = Q3 - Q1
      print(IQR)
```

```
index          43.000000
wheel-base     6.700000
length         18.200000
```

```
horsepower          55.000000
average-mileage     12.000000
price               11312.000000
price z scores       1.007998
dtype: float64
```

```
[33]: upper = Q3 + 1.5*IQR

lower = Q1 - 1.5*IQR

print(upper)

print(lower)
```

```
index              125.500000
wheel-base        111.250000
length            204.600000
horsepower         205.500000
average-mileage     49.000000
price              35088.500000
price z scores      1.755575
dtype: float64
index              -46.500000
wheel-base         84.450000
length             131.800000
horsepower         -14.500000
average-mileage      1.000000
price             -10159.500000
price z scores      -2.276416
dtype: float64
```

```
[34]: Q1 = df['price'].quantile(0.25)
Q3 = df['price'].quantile(0.75)
IQR = Q3 - Q1

upper = Q3 + 1.5*IQR

lower = Q1 - 1.5*IQR

print(upper)
print(lower)
```

```
35088.5
-10159.5
```

```
[35]: #print outliers
outliers = df['price'][((df['price'] < (Q1 - 1.5 * IQR)) | (df['price'] > (Q3 + 1.5 * IQR)))]
print(outliers)
```

```
11    41315.0
12    36880.0
25    35550.0
26    36000.0
34    40960.0
35    45400.0
46    37028.0
Name: price, dtype: float64
```

```
[36]: #replace some of the outlier values - approach 1

df['price'].replace([41315.0, 36880.0], [111, 111], inplace=True)

outliers = df['price'][((df['price'] < (Q1 - 1.5 * IQR)) | (df['price'] > (Q3 + 1.5 * IQR)))]
print(outliers)
```

```
25    35550.0
26    36000.0
34    40960.0
35    45400.0
46    37028.0
Name: price, dtype: float64
```

```
[17]: #drop the row that contains the outlier - approach 2
df.drop(index=34, inplace = True)

print("New Shape: ", df.shape)
```

```
New Shape: (60, 11)
```

```
[37]: outliers = df['price'][((df['price'] < (Q1 - 1.5 * IQR)) | (df['price'] > (Q3 + 1.5 * IQR)))]
print(outliers)
```

```
25    35550.0
26    36000.0
34    40960.0
35    45400.0
46    37028.0
Name: price, dtype: float64
```

[]: